

SUJAN REDDY

Sunnyvale, California

☎ +1 (716) 750-8189 🌐 sujanreddy.dev ✉ sujanreddyp.dev@gmail.com

Education

University at Buffalo

Master of Science in Computer Science

Aug. 2024 – May 2026

Buffalo, NY

Global Academy of Technology

Bachelor of Technology in Computer Science & Engineering

June 2020 – June 2024

Bangalore, India

Technical Skills

Languages: Python, SQL, JavaScript, Java

AI/ML: LangGraph, LangChain, LLM APIs, TensorFlow, scikit-learn

Data & APIs: PostgreSQL, SQLite, SQLAlchemy, FastAPI, REST APIs

Data Engineering: PySpark, Snowflake, Docker, AWS

Methods: NLP, embeddings, vector databases, statistical modeling

Professional Experience

IBM

Aug. 2023 – Feb. 2024

Software Engineer Assistantship

Bangalore, India

- Worked with the team under guidance on a Python predictive-analytics platform using TensorFlow and scikit-learn, contributing to a 25% improvement in prediction accuracy across 10,000 patient records.
- Supported data scraping and API automation for patient datasets, improving scraping efficiency by 98% and reducing manual data entry by 90%.
- Incorporated team feedback while helping improve the solution from data preparation through API-based retrieval.

Projects

FedLens – Agentic Federal Spending Analyst | *Python, FastAPI, LangGraph, LangChain, PostgreSQL*

GitHub

- Built an agentic workflow that converts natural-language spending questions into typed queries, retrieves verified data from PostgreSQL, and returns source-backed explanations and visual summaries.
- Implemented LangGraph orchestration, USAspending.gov ingestion, deterministic analysis, provenance metadata, and focused tests behind a FastAPI backend.

Convergence – Polymarket Research and Signal System | *Python, asyncio, FastAPI, Pandas, SQLite*

GitHub

- Built asynchronous collectors for Polymarket APIs, joined trade activity with resolved outcomes, and persisted reproducible analytical artifacts in Parquet, SQLite, and structured JSON.
- Implemented deterministic data-quality filters, statistical scoring, Bayesian aggregation, live WebSocket consensus detection, and read-only FastAPI endpoints for auditable signals.

Local Lore – Map-First San Francisco Discovery Experience | *TypeScript, React, Next.js, MapLibre, GeoJSON*

GitHub

- Built and deployed a map-first discovery platform whose local catalog agent maintains Git-tracked content, checks approved sources, validates and deduplicates facts, and publishes valid updates through GitHub and Vercel.
- Implemented MapLibre maps, local PMTiles data, geographic activity zones, clustering, responsive controls, source attribution, date filtering, accessibility context, and reduced-motion support.

Achievements

- 1st place (Bangalore) in IBM Z Datathon (International; 3,000+ participants) for Multi-level Health Analysis.
- 1st place in Hack-A-League (National; 200+ participants) for Semantic Analyzer using NLP.